

Model Optimization and Tuning Phase Report

| | |
|---------------|--|
| Date | 23 September 2024 |
| Team ID | LTVIP2024TMID24986 |
| Project Title | Movie Box Office Gross Prediction using Machine Learning |
| Maximum Marks | 10 Marks |

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (6 Marks):

| Model | Tuned Hyperparameters | Optimal Values |
|-------------------|--|--|
| Linear Regression | <pre>from sklearn.linear_model import LinearRegression from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score def linear_regression(X_train, X_test, y_train, y_test): lr = LinearRegression() lr.fit(X_train, y_train) y_pred = lr.predict(X_test) print("***Linear Regression***") print("Mean Absolute Error: ", mean_absolute_error(y_test, y_pred)) print("Mean Squared Error: ", mean_squared_error(y_test, y_pred)) print("R2 Score: ", r2_score(y_test, y_pred))</pre> | <p>Suggested code may be subject to a licence varchanaier/weather_app_vwcode</p> <pre>linear_regression(x_train, x_test, y_train, y_test)</pre> <p>**Linear Regression** Mean Absolute Error: 54.37754333076527 Mean Squared Error: 8649.147282950133 R2 Score: 0.7758003459046133</p> |
| Ridge regression | <pre>def ridge_regression_with_tuning(X_train, X_test, y_train, y_test): param_grid = {'alpha': [0.001, 0.01, 0.1, 1, 10, 100]} # Regularization strength ridge = Ridge() grid_search = GridSearchCV(estimator=ridge, param_grid=param_grid, cv=5, scoring='r2') grid_search.fit(X_train, y_train) best_ridge = grid_search.best_estimator_ y_pred = best_ridge.predict(X_test) print("Best Ridge Hyperparameters:", grid_search.best_params_)</pre> | <pre>ridge_regression_with_tuning(x_train, x_test, y_train, y_test)</pre> <p>Best Ridge Hyperparameters: {'alpha': 100} R2 Score: 0.7767821907211369 Mean Absolute Error: 54.220345539504514 Mean Squared Error: 8611.26979174089</p> |

| | | |
|-----------------------|--|--|
| SVM regression | <pre>def svr_with_tuning(X_train, X_test, y_train, y_test): param_grid = { 'C': [0.1, 1, 10, 100], # Regularization parameter 'epsilon': [0.001, 0.01, 0.1, 1], # Epsilon-tube within which no penalty is assessed 'kernel': ['linear', 'rbf'] # Type of kernel } svr = SVR() grid_search = GridSearchCV(estimator=svr, param_grid=param_grid, cv=5, scoring='r2') grid_search.fit(X_train, y_train) best_svr = grid_search.best_estimator_</pre> | <p>Best SVR Hyperparameters: {'C': 100, 'epsilon': 0.1, 'kernel': 'linear'}</p> <p>R2 Score: 0.7367468556578113</p> <p>Mean Absolute Error: 52.433519903226696</p> <p>Mean Squared Error: 10155.748131291039</p> |
| Random Forest | <pre>def random_forest_regressor_with_tuning(X_train, X_test, y_train, y_test): param_grid = { 'n_estimators': [100, 200, 300], # Number of trees in the forest 'max_depth': [10, 20, 30, None], # Maximum depth of the tree 'min_samples_split': [2, 5, 10], # Minimum number of samples required to split 'min_samples_leaf': [1, 2, 4], # Minimum number of samples required at a leaf 'bootstrap': [True, False] # Whether bootstrap samples are used when } rf = RandomForestRegressor() grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='r2') grid_search.fit(X_train, y_train) best_rf = grid_search.best_estimator_ y_pred = best_rf.predict(X_test)</pre> | <p>Best Random Forest Hyperparameters: {'bootstrap': True, 'max_depth': 30, 'min_samples_split': 5, 'min_samples_leaf': 4, 'n_estimators': 200}</p> <p>R2 Score: 0.7480565218845108</p> <p>Mean Absolute Error: 51.16163391161018</p> <p>Mean Squared Error: 9719.445188227117</p> |
| Decision Tree | <pre>def decision_tree_regressor_with_tuning(X_train, X_test, y_train, y_test): dt = DecisionTreeRegressor() param_grid = { 'max_depth': [5, 10, 20], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] } grid_search = GridSearchCV(dt, param_grid, cv=5, scoring='r2') grid_search.fit(X_train, y_train) best_dt = grid_search.best_estimator_ y_pred = best_dt.predict(X_test)</pre> | <p>**Decision Tree Regressor**</p> <p>Best Params: {'max_depth': 5, 'min_samples_leaf': 4, 'min_samples_split': 5}</p> <p>Mean Absolute Error: 55.310368452997785</p> <p>Mean Squared Error: 10950.050909112899</p> <p>R2 Score: 0.7161572643132731</p> |

Performance Metrics Comparison Report (2 Marks):

| Model | Baseline Metric | Optimized Metric |
|--------------------------|-----------------|------------------|
| Linear Regression | R2: 0.7758 | R2: 0.7758 |
| SVM Regression | R2: 0.1528 | R2: 0.1528 |
| Random Forest Regression | R2: 0.7617 | R2: 0.7617 |
| Decision Tree Regressor | R2: 0.5318 | R2: 0.5318 |

Final Model Selection Justification (2 Marks):

| Final Model | Reasoning |
|-------------------|--|
| Linear Regression | The Linear Regression model was chosen as the final optimized model because it exhibited the highest R-squared value (0.7758), indicating a strong fit to the data. Additionally, it had a lower MSE (8649.14) and Accuracy (77%) compared to other models, suggesting superior predictive accuracy. |